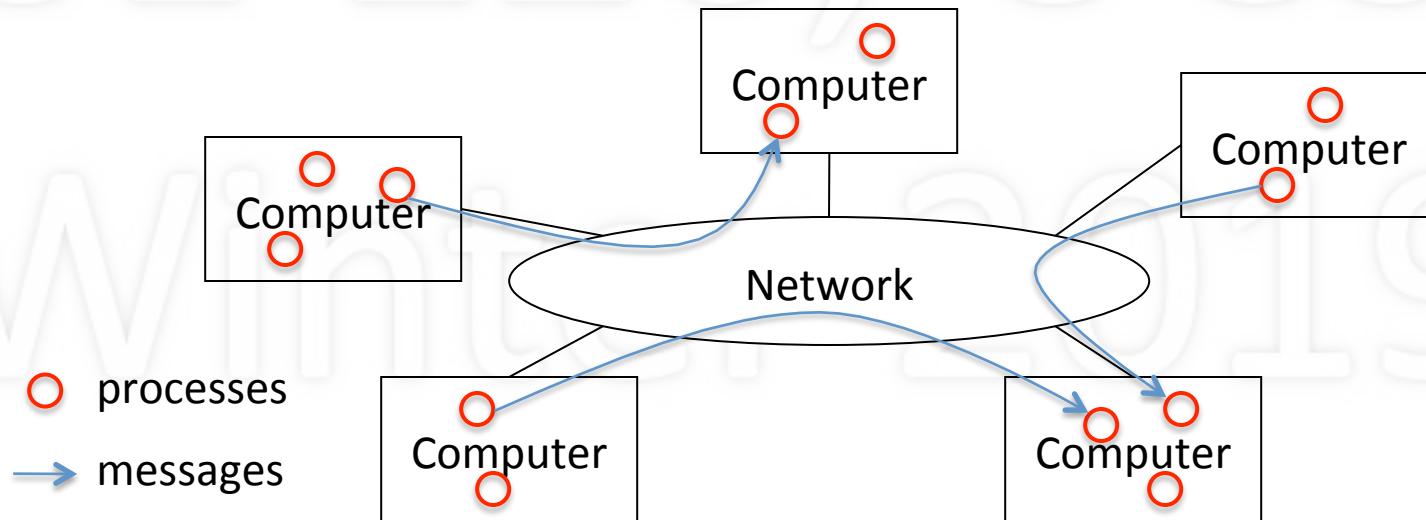


# CSE 120: Principles of Operating Systems

## Lecture 16: Distributed Systems

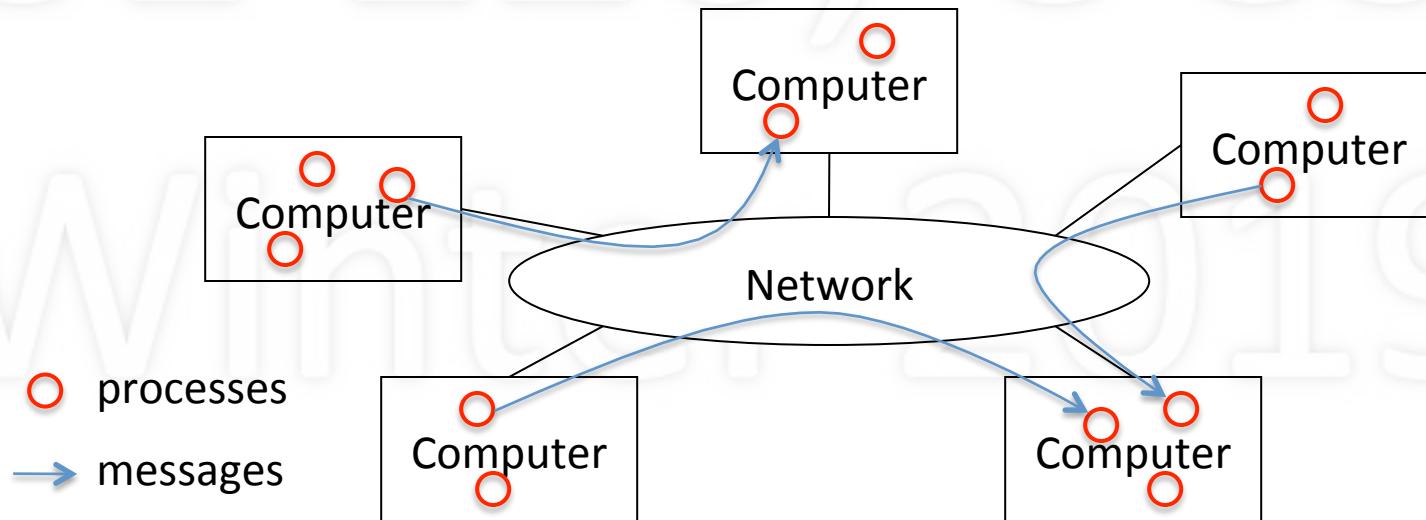
Prof. Joseph Pasquale  
University of California, San Diego  
March 13, 2019

# What is a Distributed System?



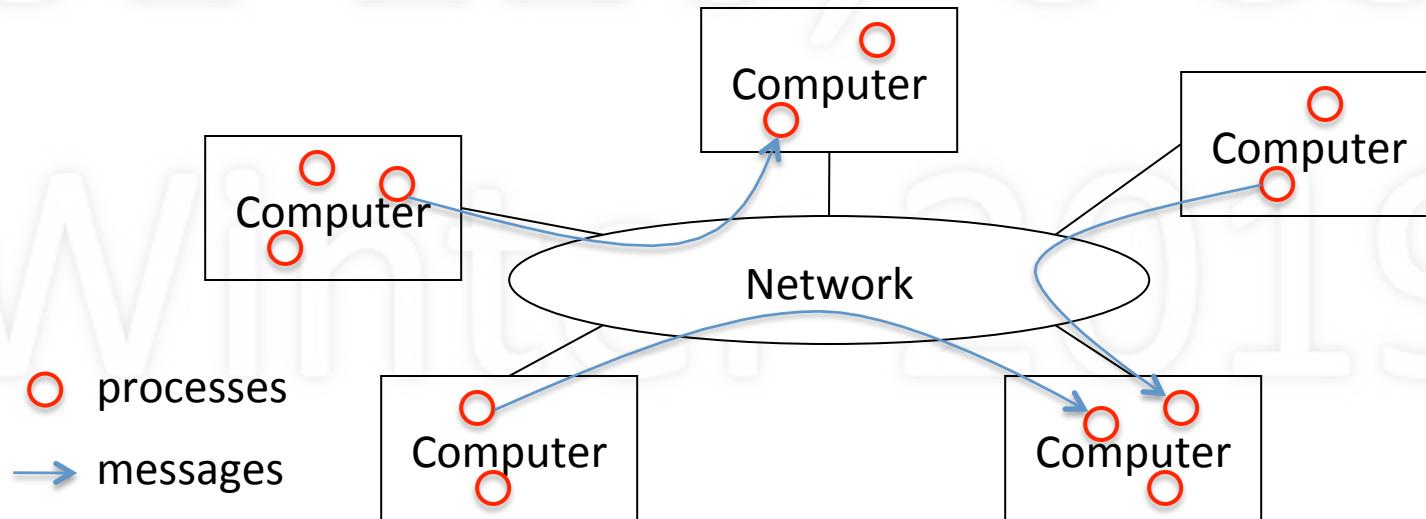
- Cooperating processes in a computer network
- Degree of integration
  - Loose: Internet applications, email, web browsing
  - Medium: remote execution, remote file systems
  - Tight: process migration, distributed file systems

# Advantages



- Speed: parallelism, less contention
- Reliability: redundancy, fault tolerance, “NSPF”
- Scalability: incremental growth, economy of scale
- Geographic distribution: low latency, reliability

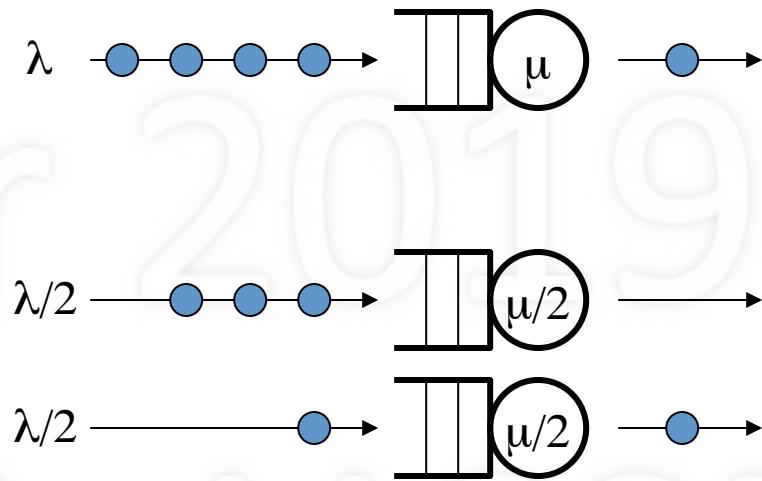
# Disadvantages



- Fundamental problems of decentralized control
  - State uncertainty: no shared memory or clock
  - Action uncertainty: mutually conflicting decisions
- Distributed algorithms are complex

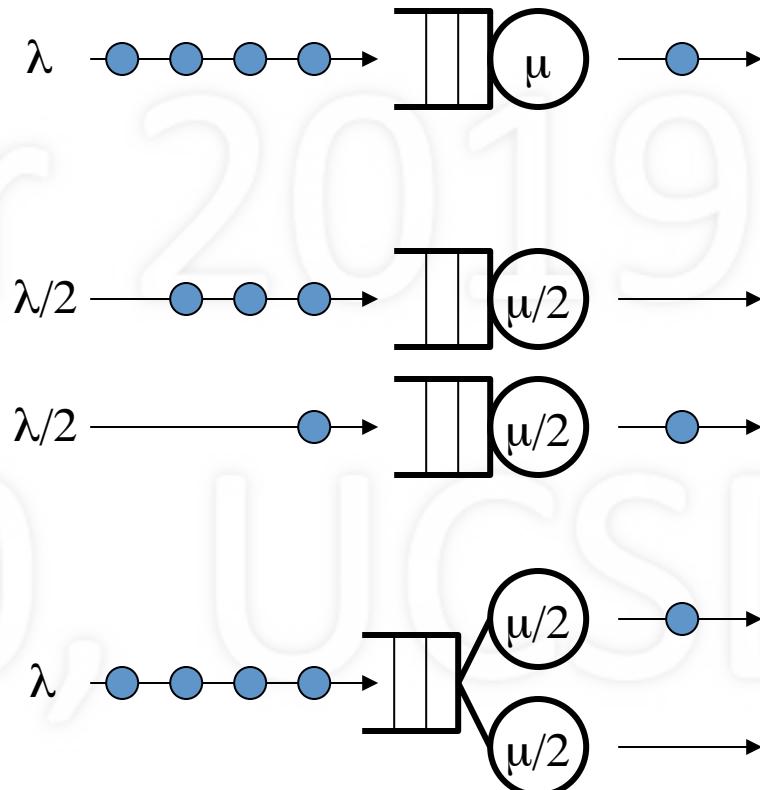
# Is Distribution Better?

- Single fast server with single queue
- Multiple slower servers with separate queues

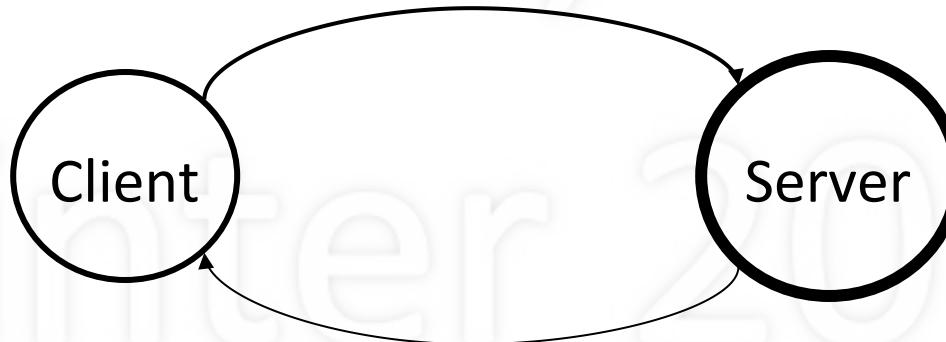


# Is Distribution Better?

- Single fast server with single queue
- Multiple slower servers with separate queues
- Multiple slower servers, single queue
- Little's Law:  $N = \lambda W$



# The Client/Server Model



- Client
  - Short-lived process that makes requests
  - “User-side” of application
- Server
  - Exports well-defined requests/response interface
  - Long-lived process that waits for requests
  - Upon receiving request, carries it out (may spawn)

# Peer-to-Peer



- A peer talks directly with another peer
  - No intermediary (e.g., central server) involved
  - Symmetric (unlike asymmetric client/server)
- In actuality, may be dynamic client/server
  - A requests file from B; A acts as client, B as server
  - C can now request file from A; A acts as server

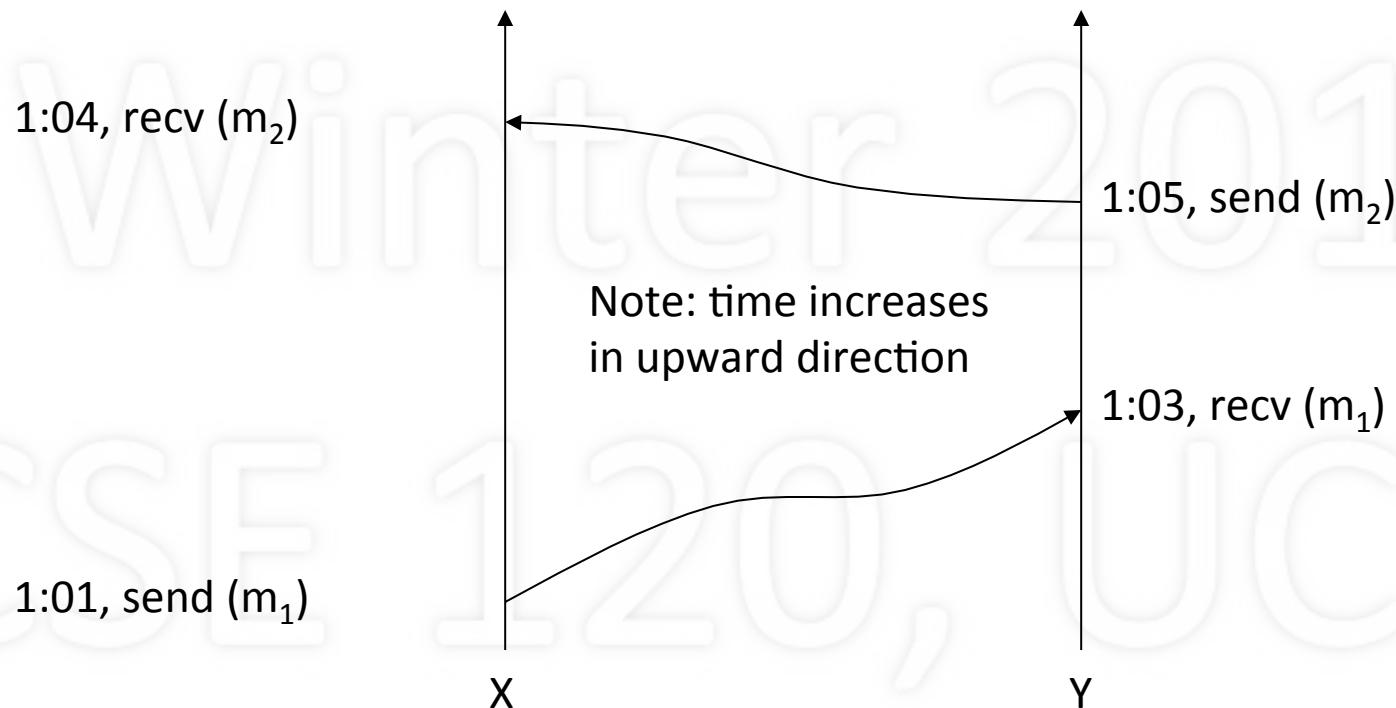
# Distributed Algorithms

- Event ordering
- Leader election
- Mutual exclusion

# Event Ordering

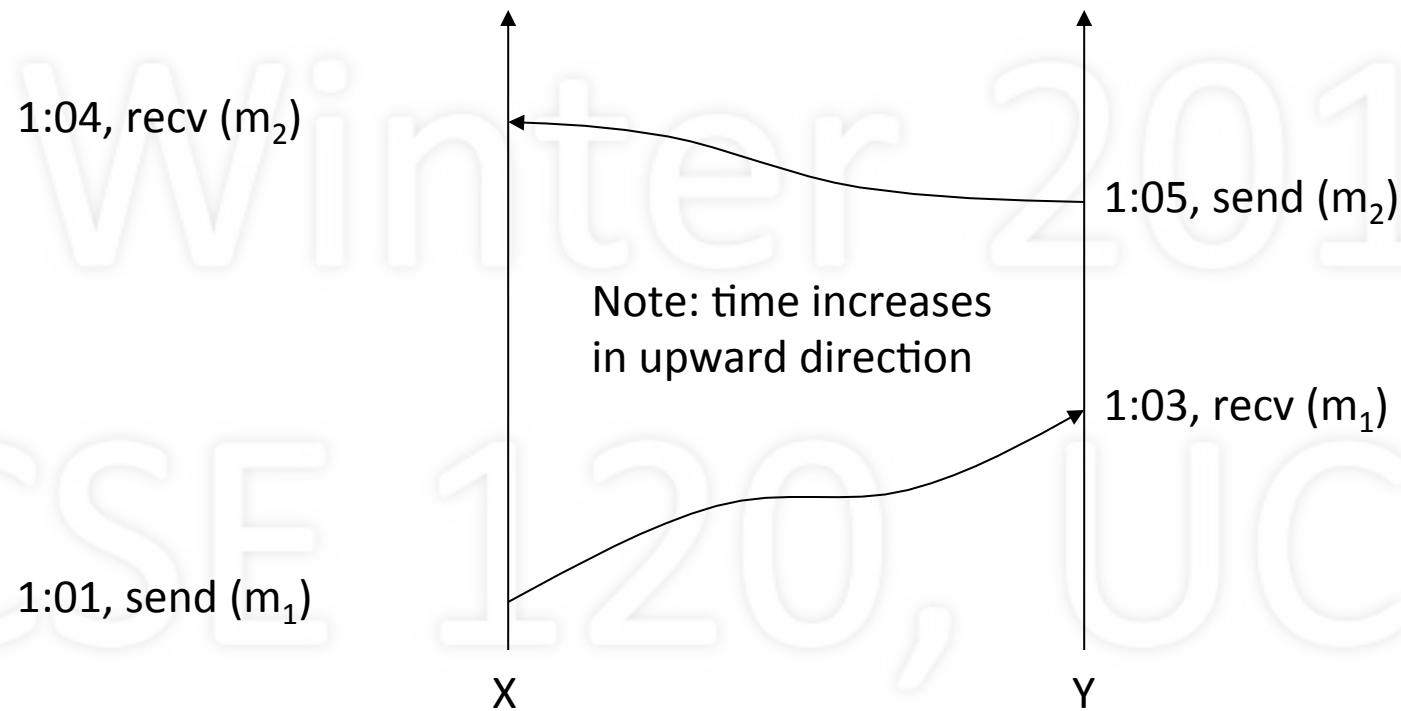
- Order events given no shared clock/memory
- Happened-before relation: →
  - A, B events in same process and A before B:  $A \rightarrow B$
  - A is a send event, B is a receive event:  $A \rightarrow B$
  - If  $A \rightarrow B$  and  $B \rightarrow C$ , then  $A \rightarrow C$
- Implementation
  - Timestamp all events based on local clock
  - Upon receiving a message, advance local clock
  - Resolve ties by ordering machines

# Event Ordering: Example



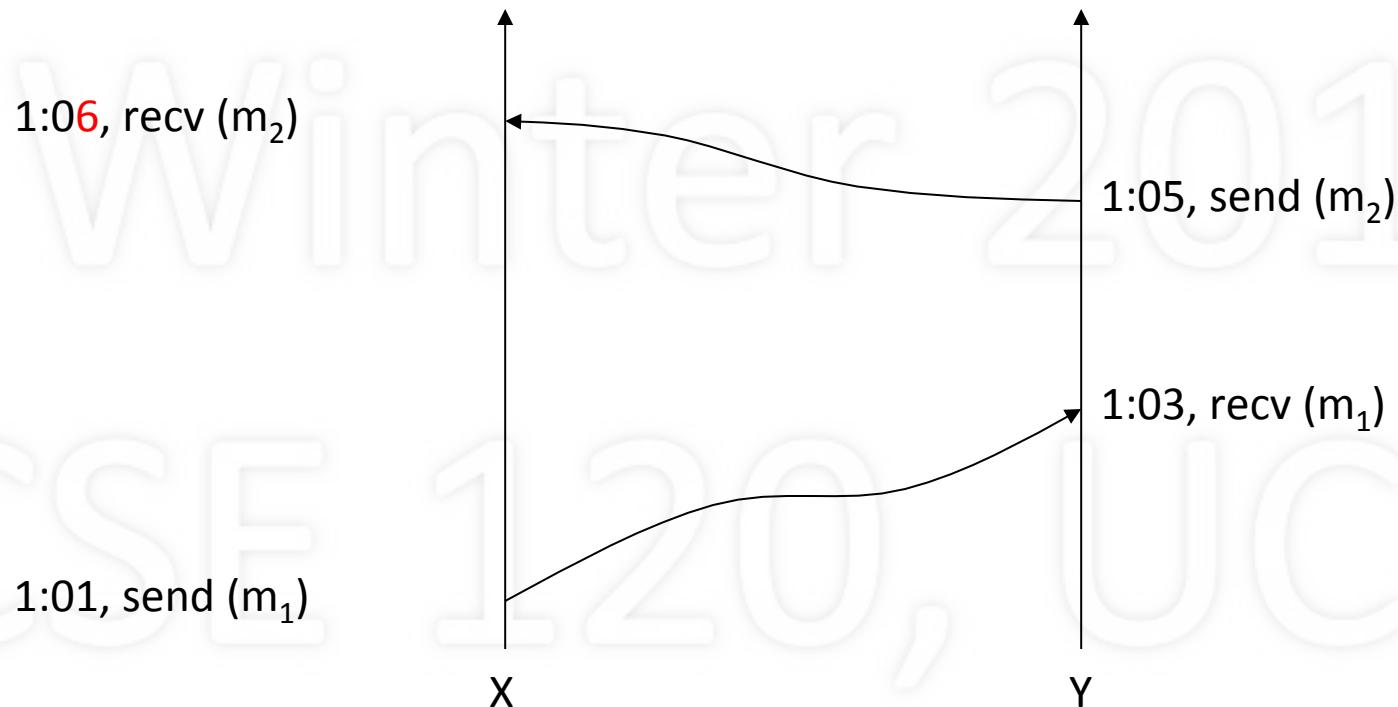
1:01 send ( $m_1$ ); 1:03 recv ( $m_1$ ); 1:05 send ( $m_2$ ); 1:04 recv ( $m_2$ );

# Event Ordering: Example



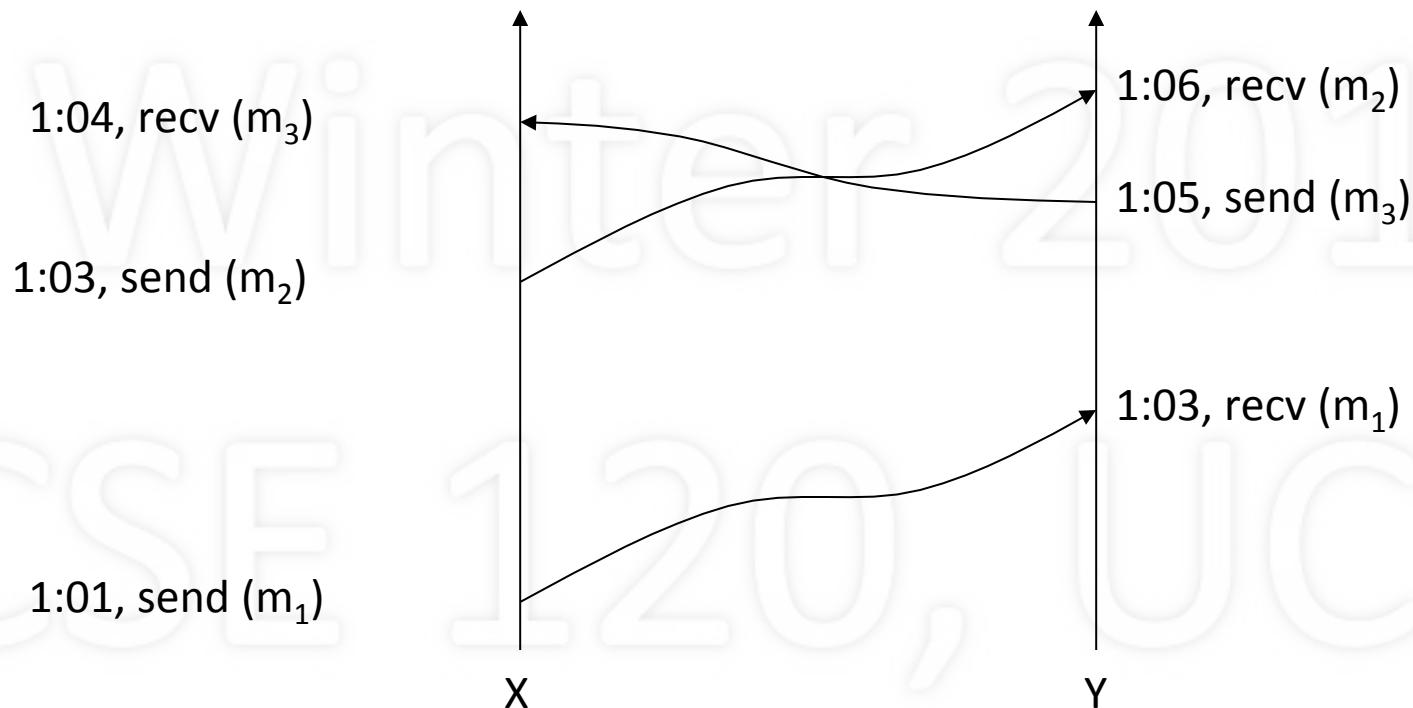
1:01 send ( $m_1$ ); 1:03 recv ( $m_1$ ); 1:05 send ( $m_2$ ); 1:04 recv ( $m_2$ );

# Event Ordering: Example

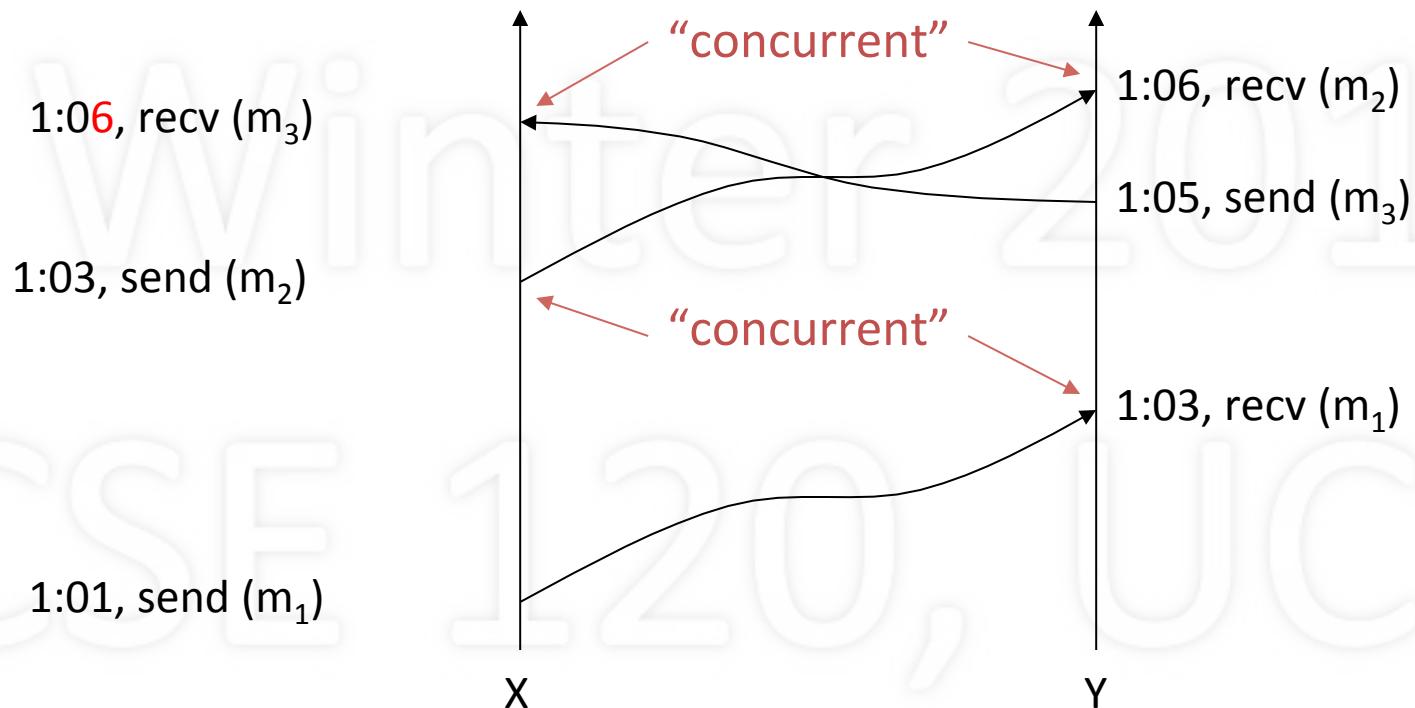


1:01 send ( $m_1$ ); 1:03 recv ( $m_1$ ); 1:05 send ( $m_2$ ); 1:06 recv ( $m_2$ )

# Event Ordering: Another Example

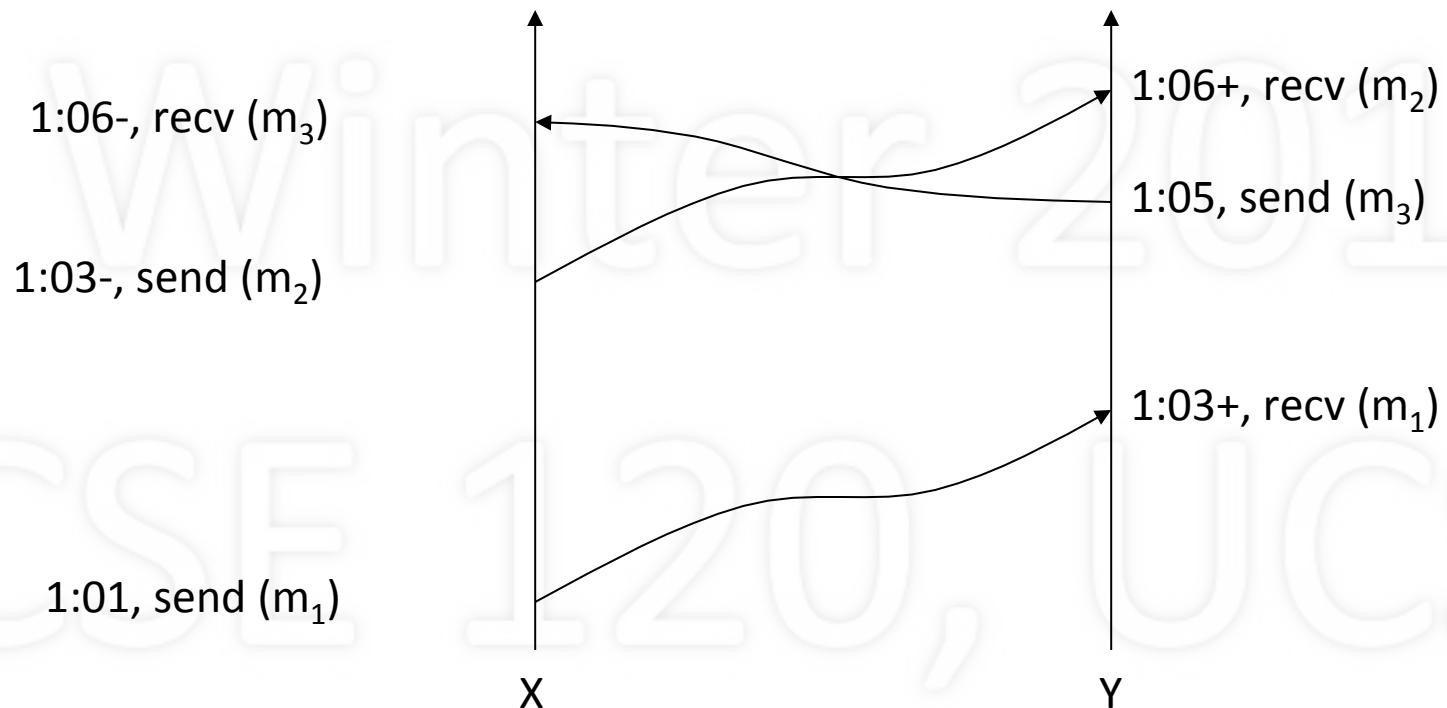


# Event Ordering: Another Example



1:01 send ( $m_1$ ); 1:03 recv ( $m_1$ ); 1:03 send ( $m_2$ ); 1:05 send ( $m_3$ ); 1:06 recv ( $m_3$ ); 1:06 recv ( $m_2$ )

# Event Ordering: Another Example

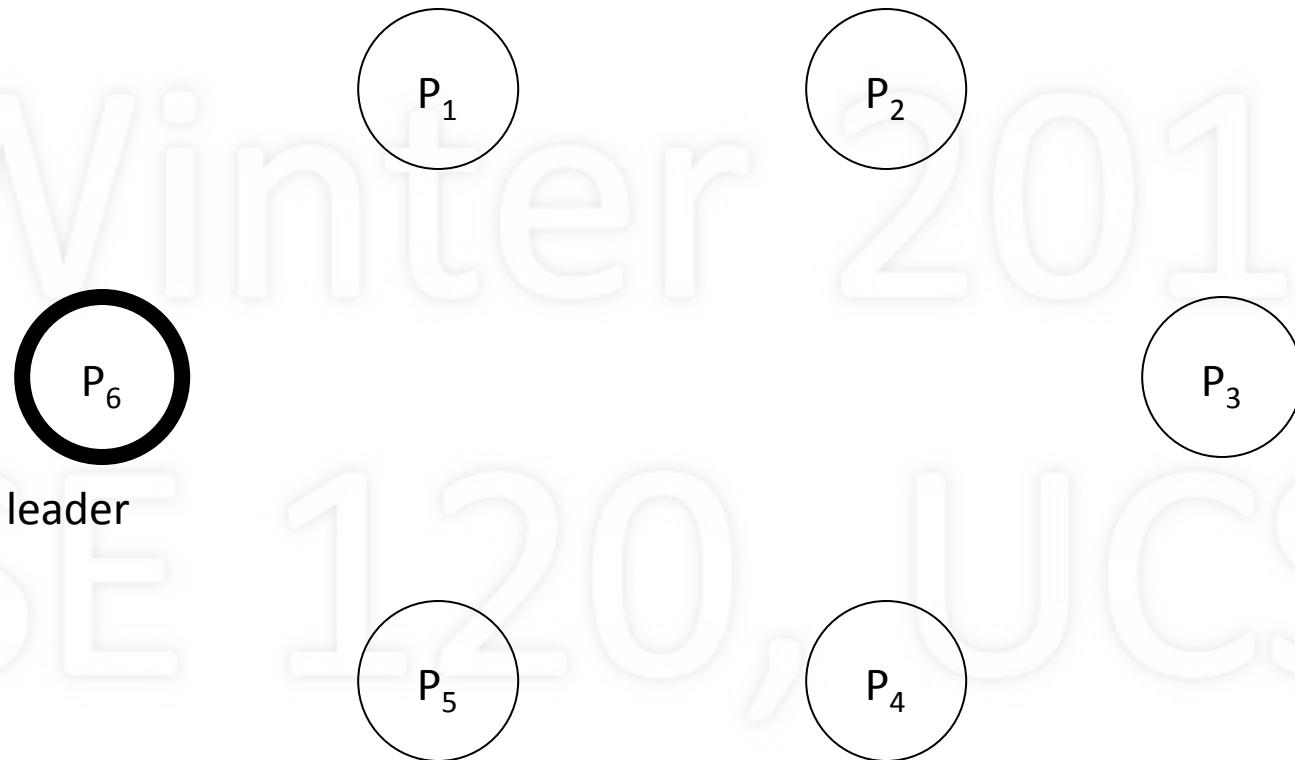


1:01 send ( $m_1$ ); 1:03- send ( $m_2$ ); 1:03+ recv ( $m_1$ ); 1:05 send ( $m_3$ ); 1:06- recv ( $m_3$ ); 1:06+ recv ( $m_2$ )

# Leader Election

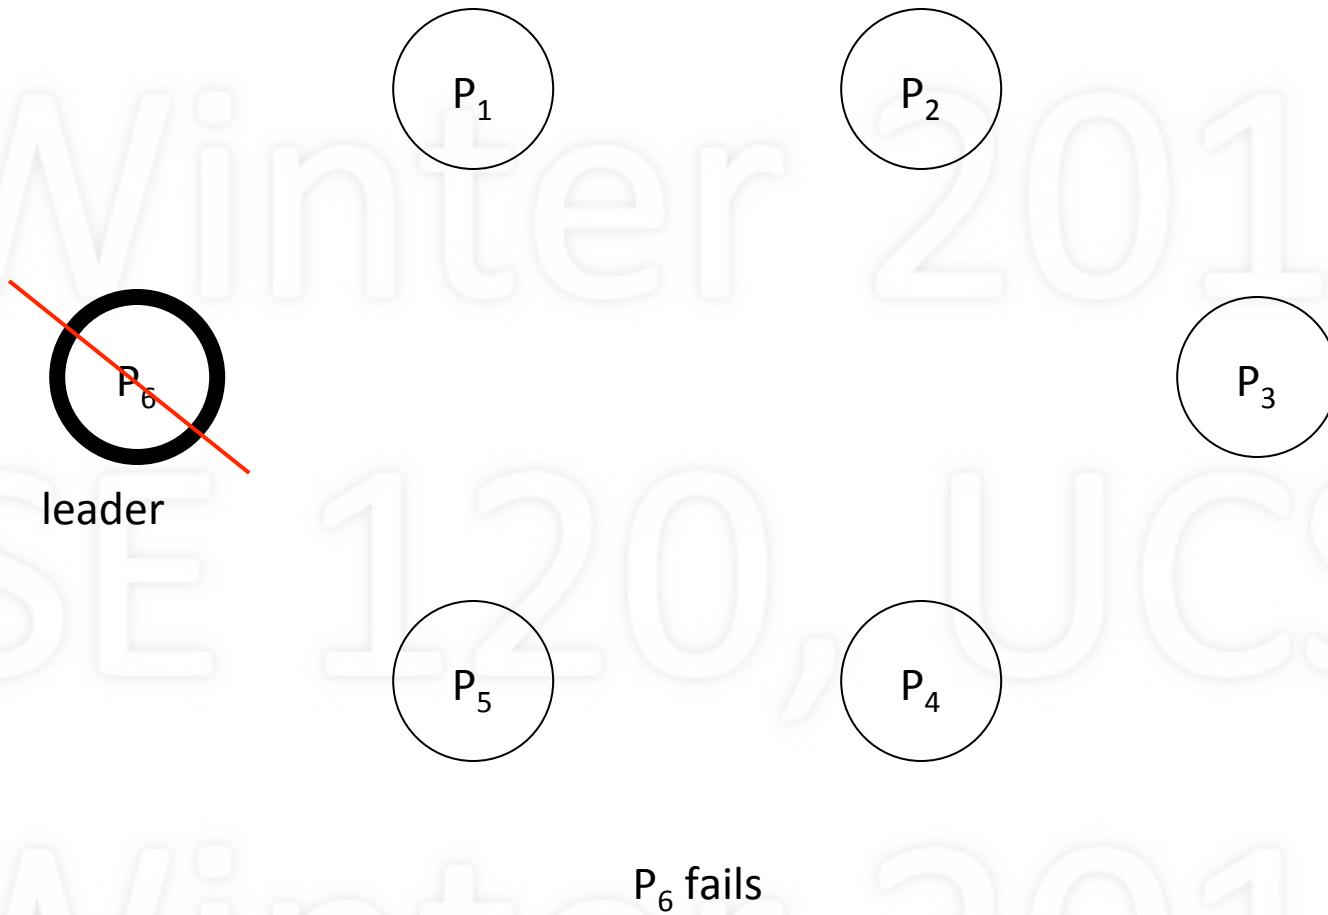
- Many distributed algorithms rely on leader
- Need to determine if leader exists; if not, elect
- Bully algorithm (elects leader L)
  - Every process is numbered (priority):  $P_1, P_2, \dots$
  - $P_j$  sends request to L, no reply; tries to elect itself
  - $P_j$  sends “Can I be leader?” to all  $P_{k>j}$
  - No replies,  $P_j$  sends to all  $P_{i<j}$ , “I am leader”, done
  - If some  $P_{k>j}$  replies,  $P_j$  lets  $P_k$  try to elect itself
  - If no message from  $P_k$ ,  $P_j$  tries to elect itself again

# Leader Election

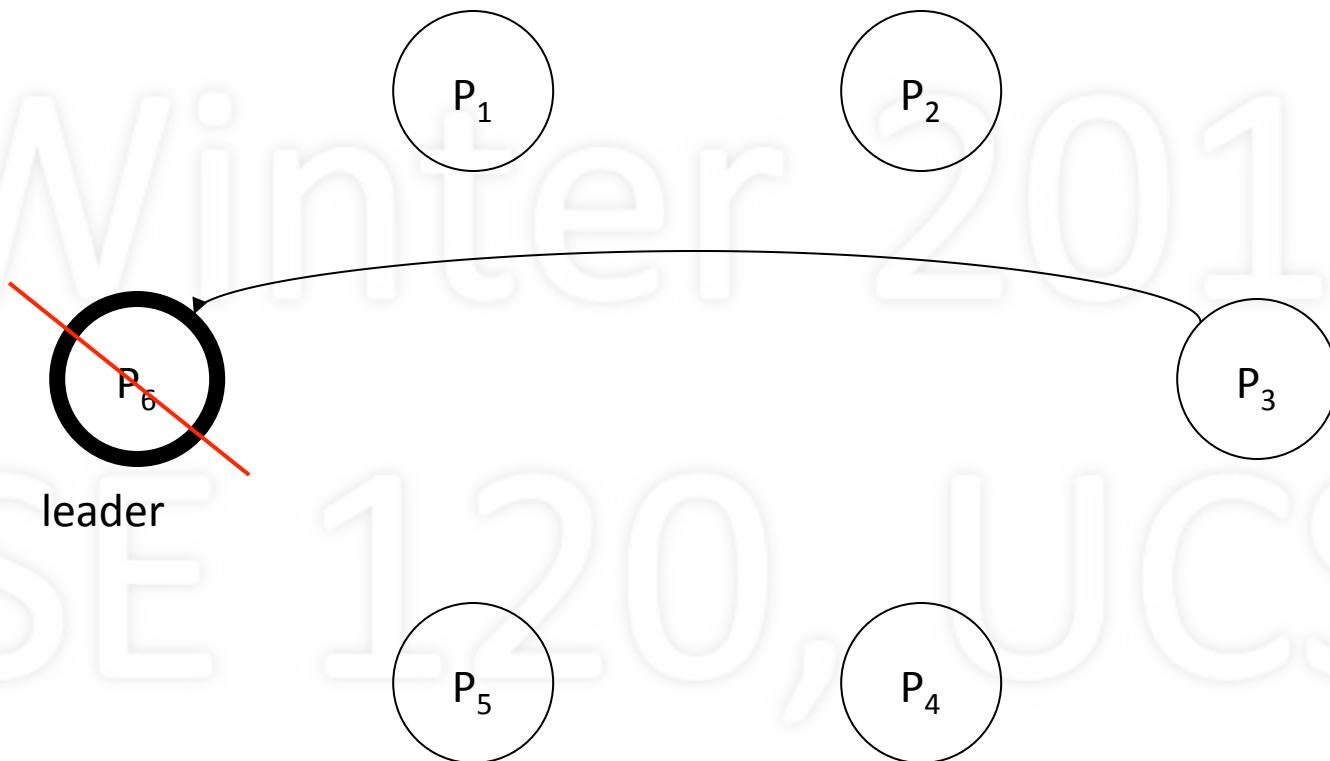


Assume  $P_6$  is leader (highest number)

# Leader Election

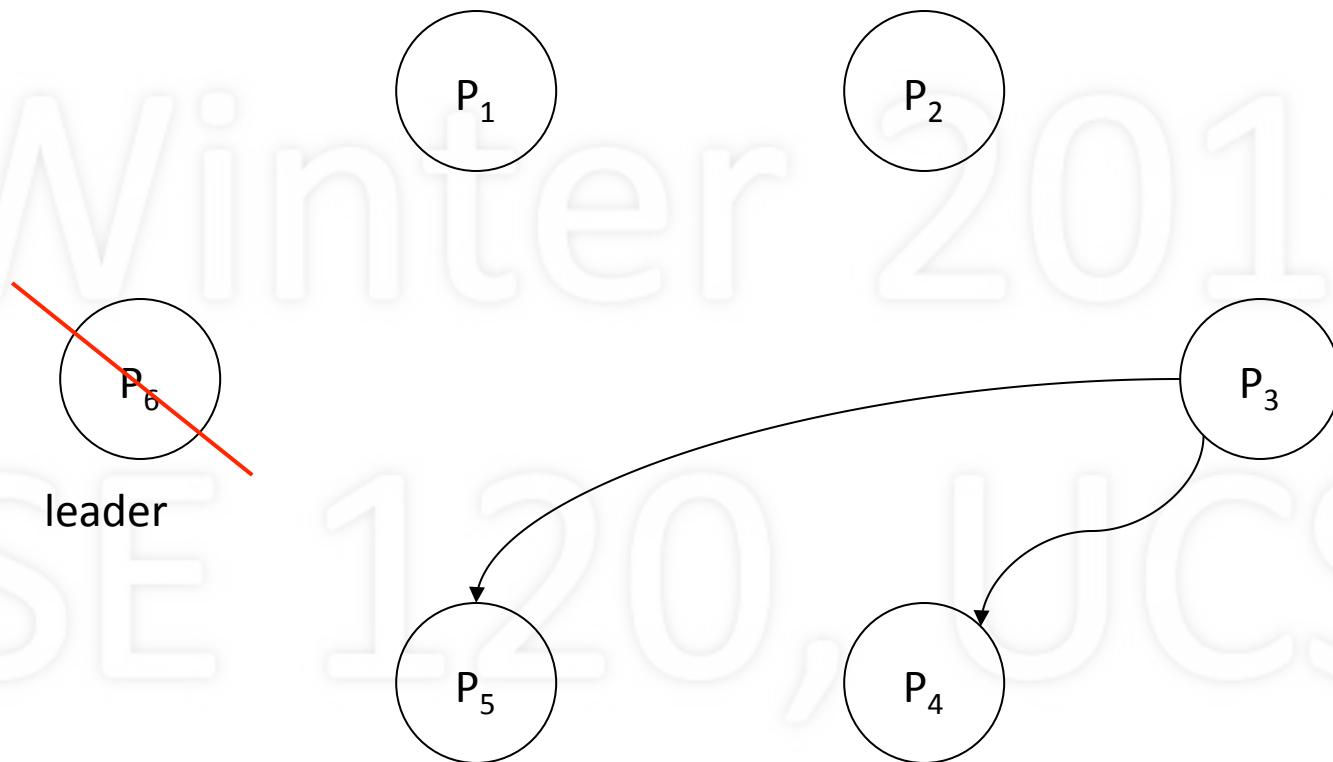


# Leader Election



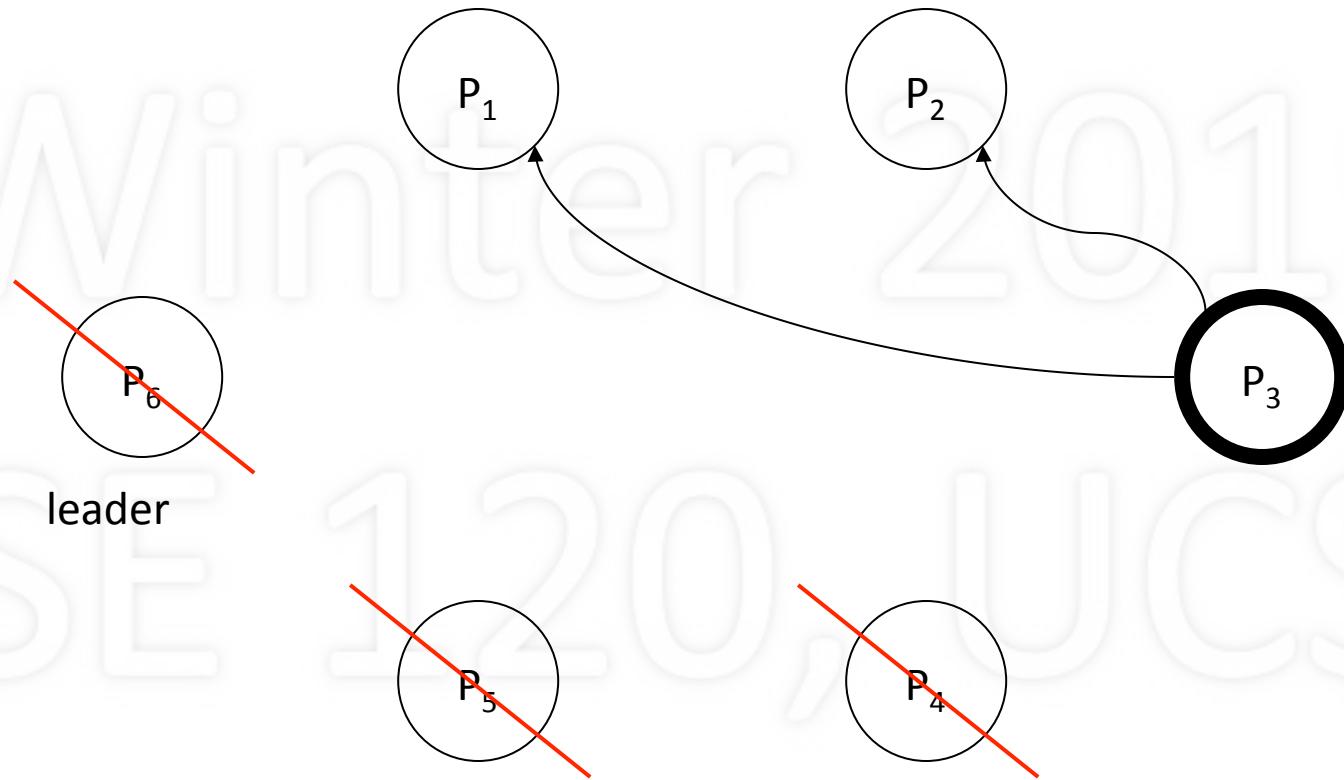
P<sub>3</sub> tries to contact leader; no reply

# Leader Election



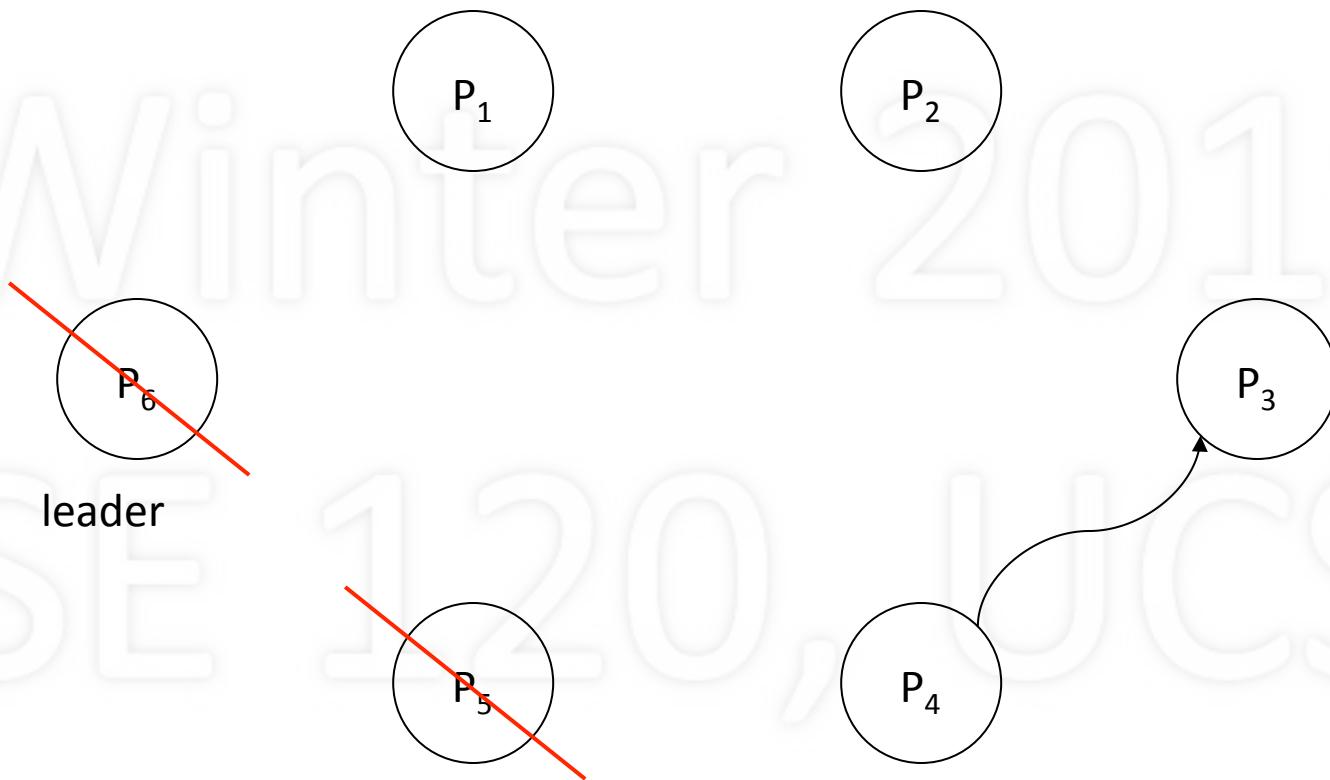
P<sub>3</sub> sends election message to P<sub>4</sub> and P<sub>5</sub>

# Leader Election



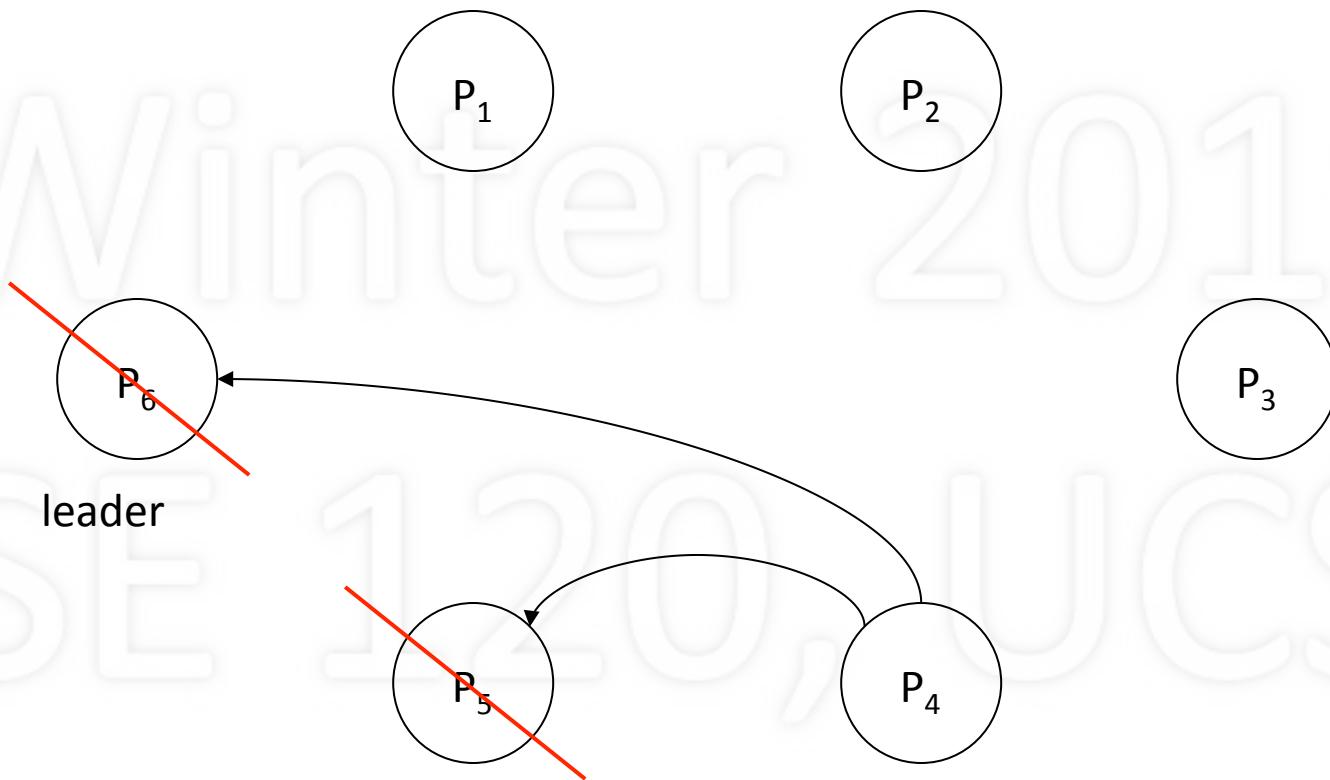
If P<sub>3</sub> were to not hear from P<sub>4</sub> and P<sub>5</sub>, it would tell P<sub>1</sub> and P<sub>2</sub> "I'm the leader"

# Leader Election



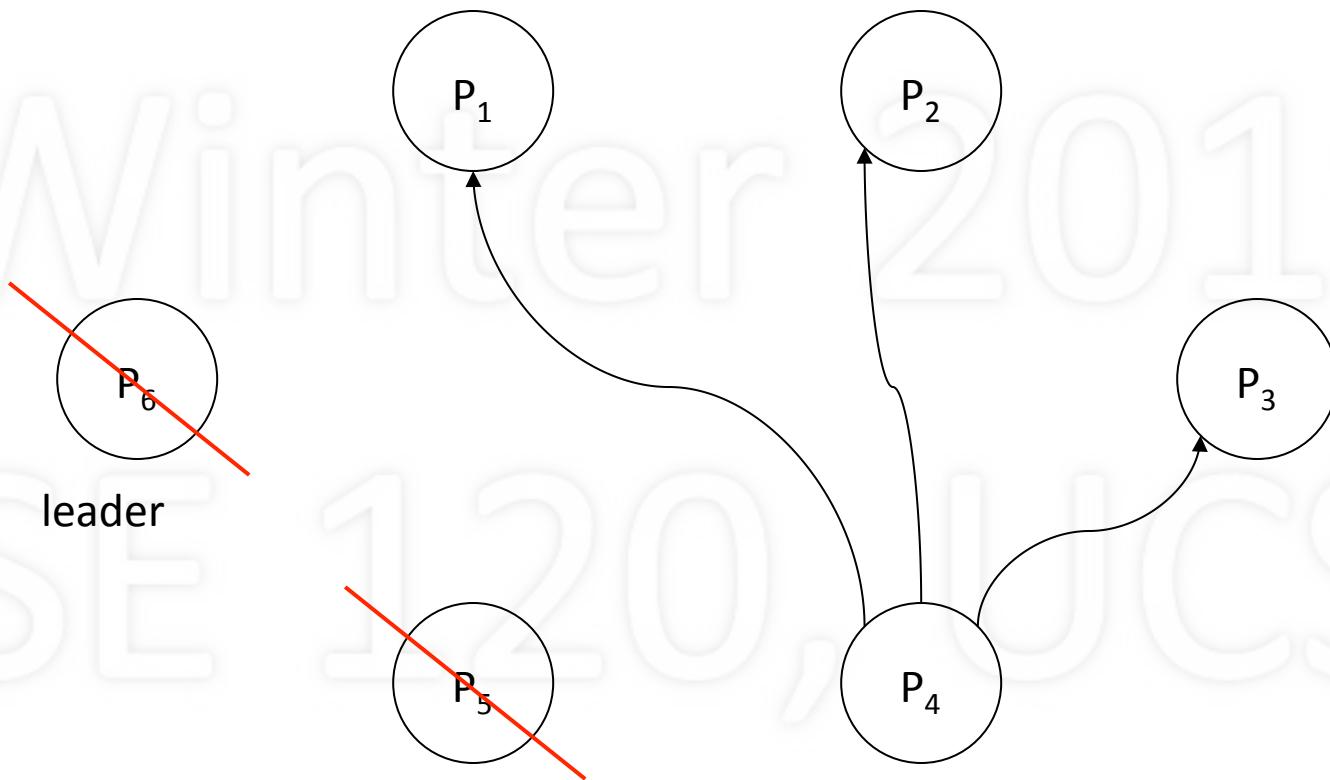
If only P<sub>4</sub> is alive, it replies to P<sub>3</sub>

# Leader Election



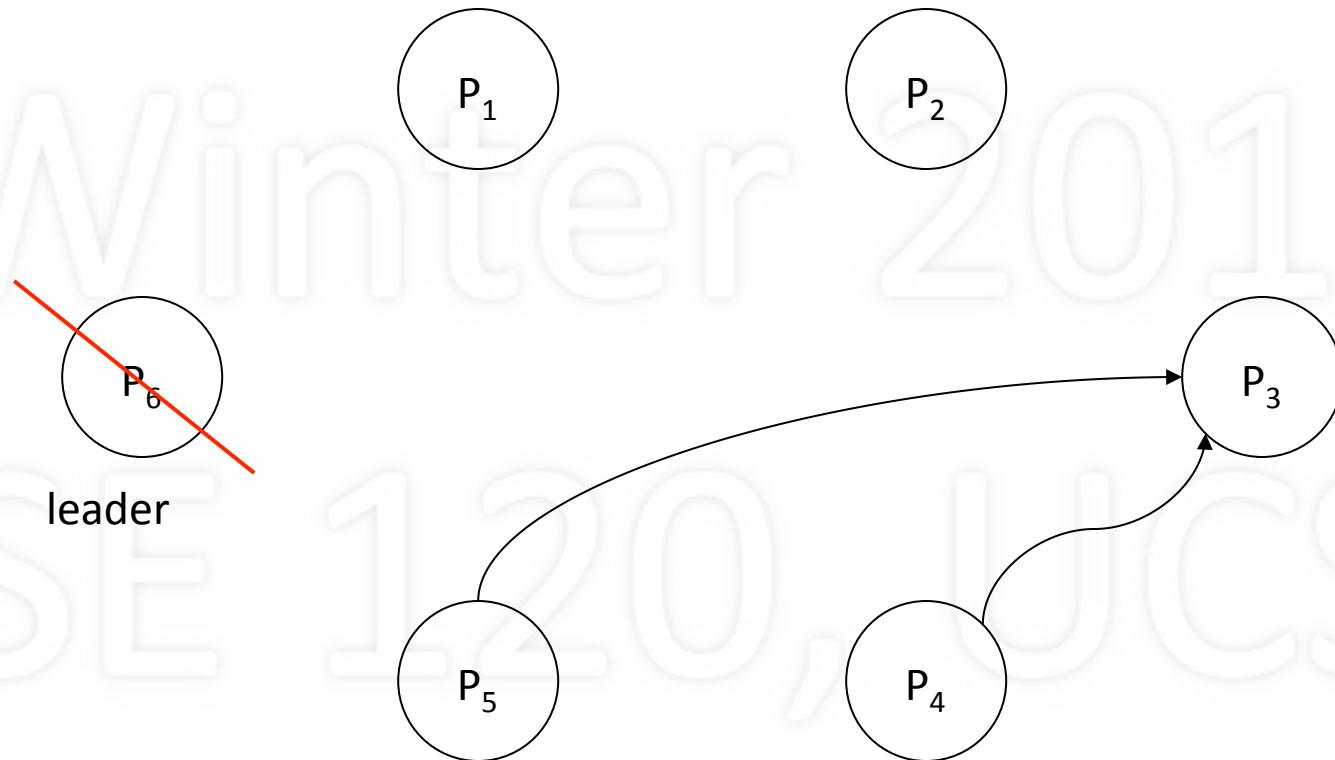
P<sub>4</sub> tries to elect itself as leader, sending messages to P<sub>5</sub> and P<sub>6</sub>; no replies

# Leader Election



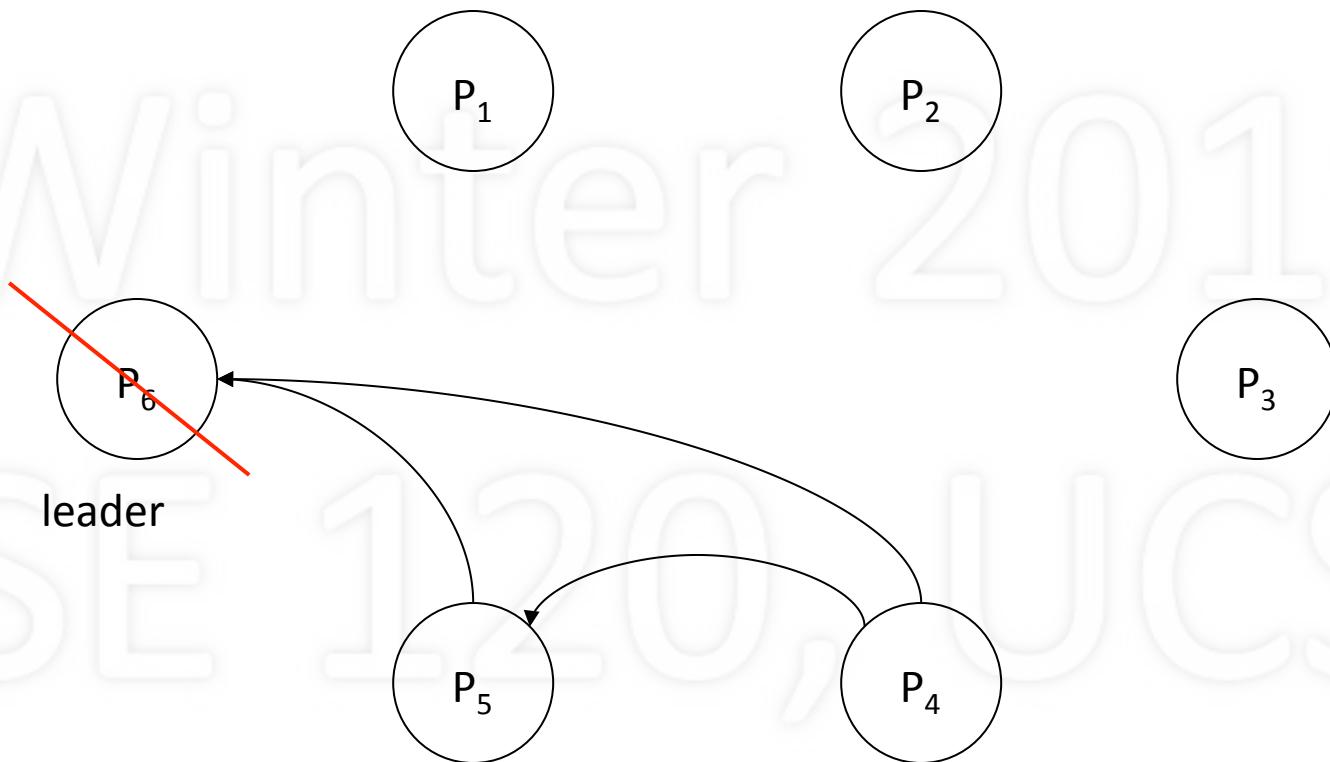
$P_4$  declares itself as leader, telling  $P_1$ ,  $P_2$ , and  $P_3$

# Leader Election



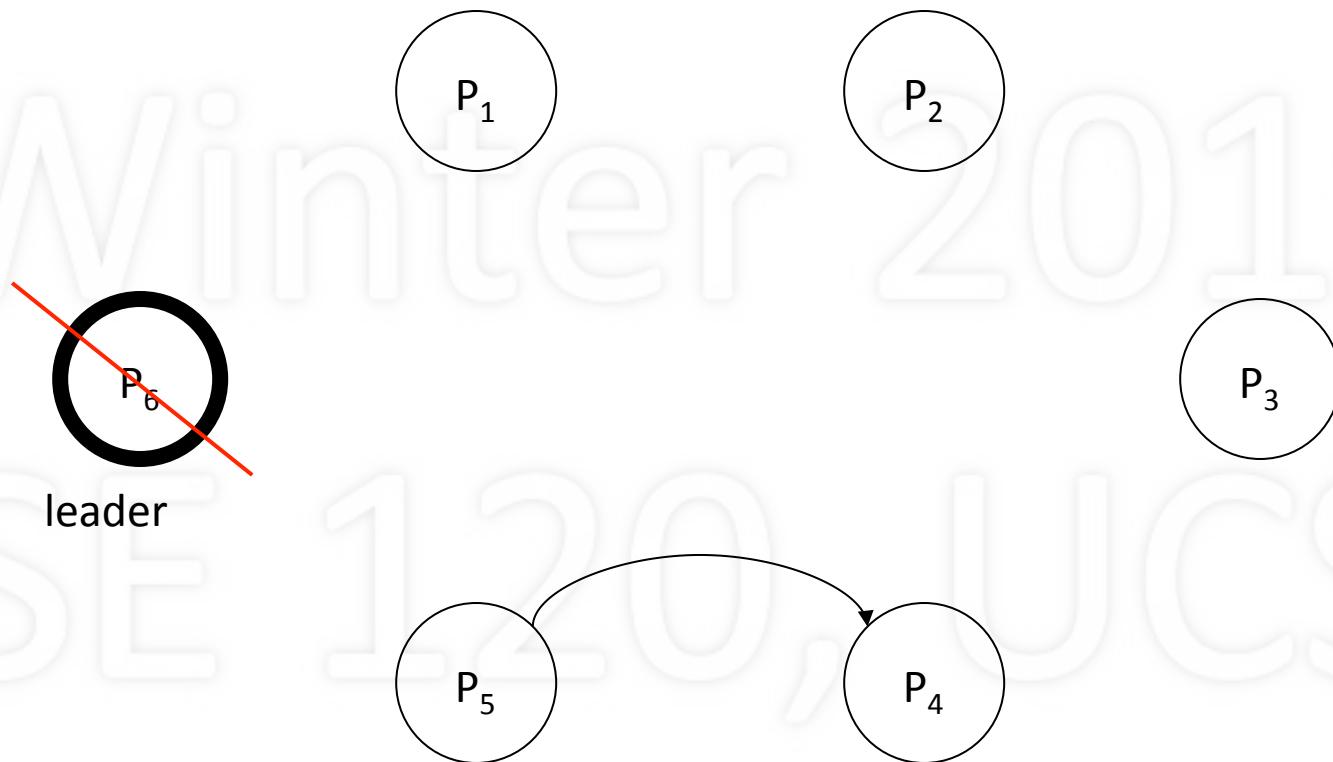
If both P<sub>4</sub> and P<sub>5</sub> are alive, they reply to P<sub>3</sub>

# Leader Election



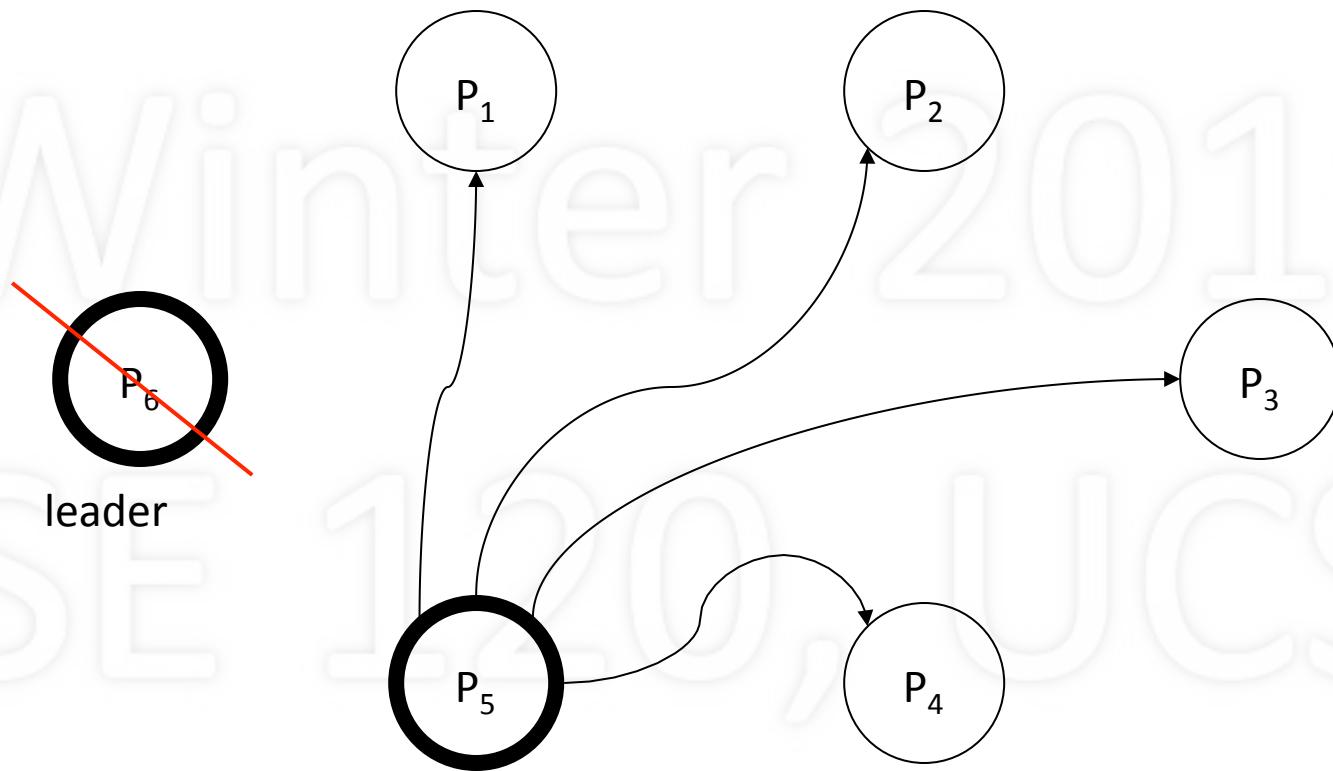
Both P<sub>4</sub> and P<sub>5</sub> try to elect themselves

# Leader Election



P<sub>4</sub> gets reply from P<sub>5</sub>, but no reply from P<sub>6</sub> to P<sub>5</sub>

# Leader Election

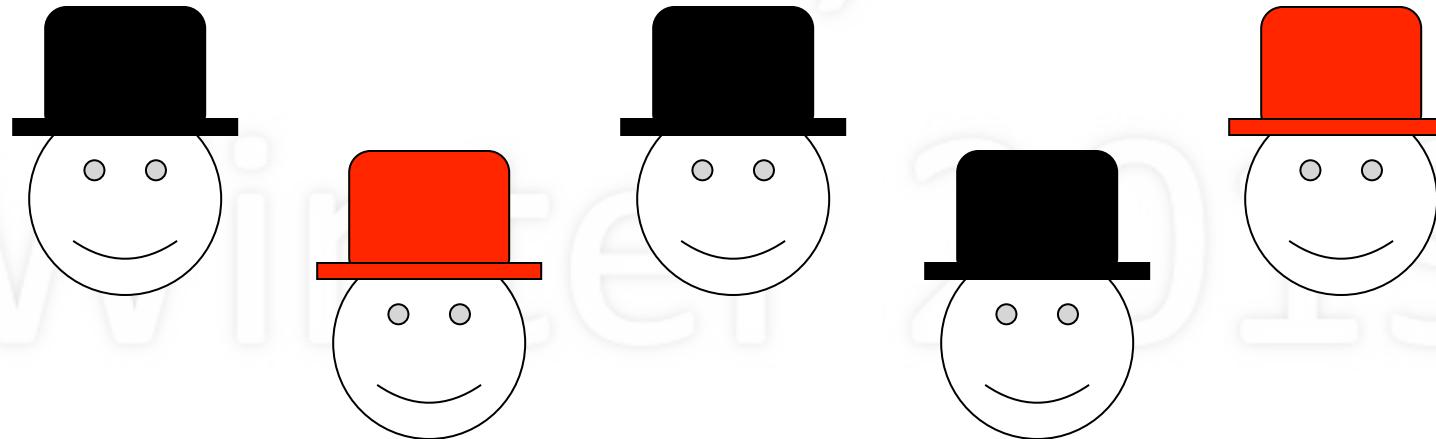


P<sub>5</sub> tells everyone below it, “I’m the leader”

# Mutual Exclusion

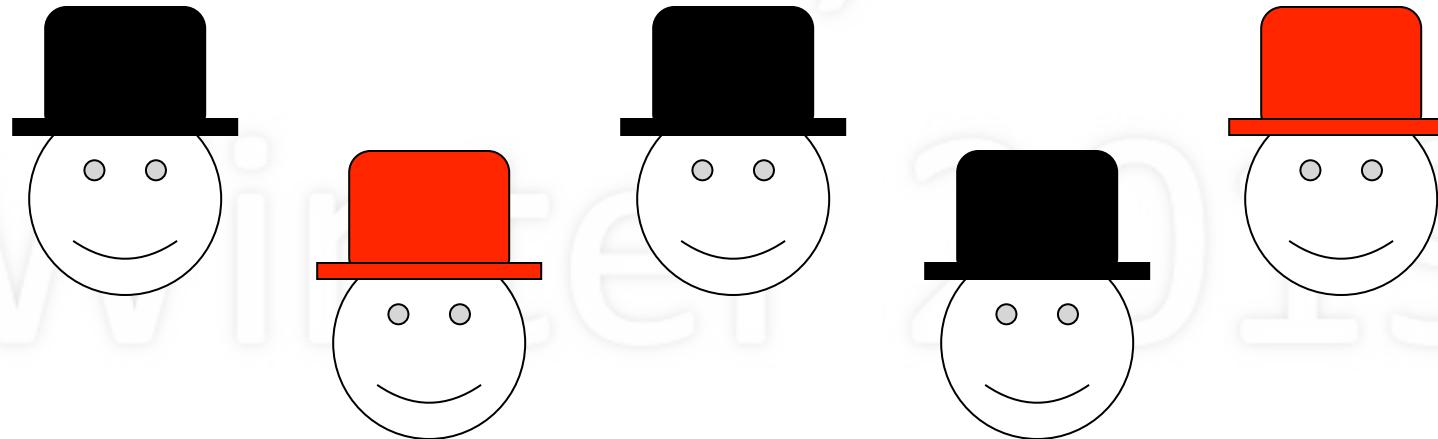
- Centralized approach
  - Single process acts as coordinator server
  - Request, reply (to allow entrance), release
- Distributed approach
  - Process sends time-stamped request to all others
  - Waits until it receives replies from all (ok to enter)
  - Enter critical section (may get requests, defers)
  - Upon exiting, responds (to release) to all deferred
  - Use timestamps to order “simultaneous” requests

# Common Knowledge Problem



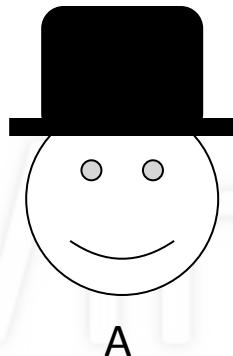
- n people, each wearing red or black hat
- Each can see the others' hats, but not own hat
- Hint: “At least one of you has a black hat”
  - Told publicly: all hear it, all know that all hear it, ...

# Common Knowledge Problem

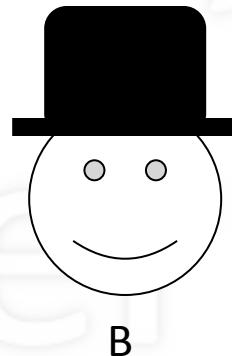


- They are asked (publicly): “Is your hat black?”
- Possible replies (made simultaneously)
  - “Yes”
  - “No”
  - “I don’t know”

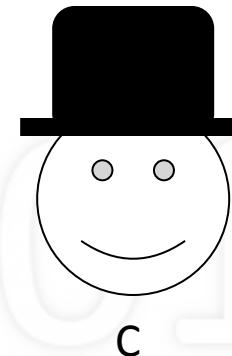
# Example: 3 Wearing Black Hats



A



B



C

All told: “At least one of you has a black hat”

Q1. “Is your hat black?”

All: “I don’t know”

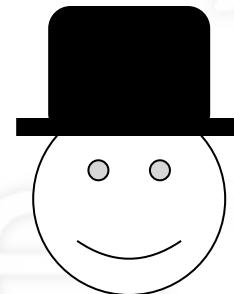
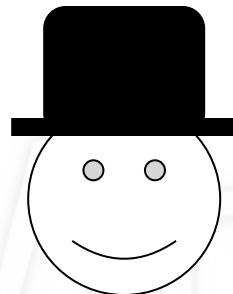
Q2. “Is your hat black?”

All: “I don’t know”

Q3. “Is your hat black?”

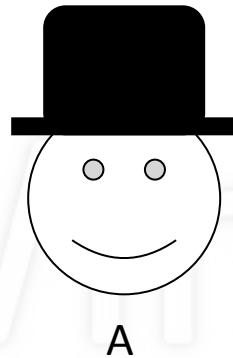
All: “Yes!”

# Questions

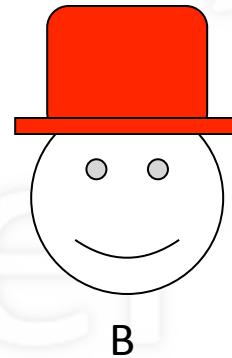


- How do each figure out they have black hat?
- Why is the hint necessary?
  - After all, each can see someone with a black hat
  - Does it have to be public?
- What is relevance to distributed systems?

# What If One Black Hat?



A



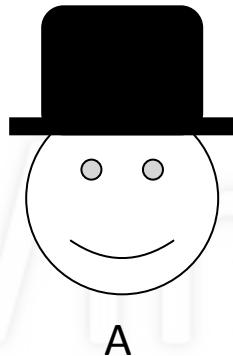
B



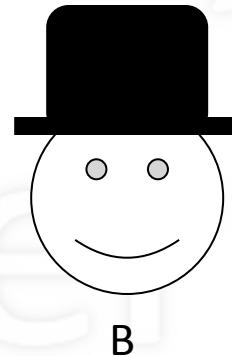
C

- A sees 2 red hats
- Can answer “Yes” after 1<sup>st</sup> question
  - Was told “at least one of you has a black hat”

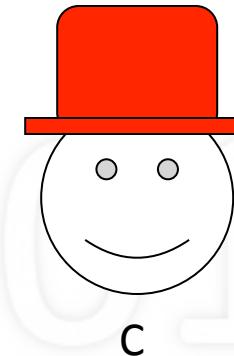
# What If Two Black Hats?



A



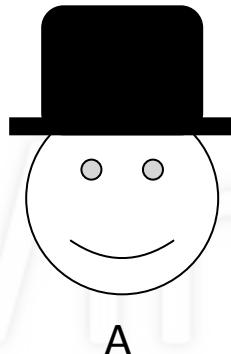
B



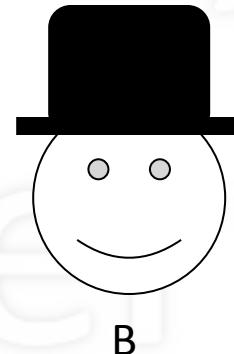
C

- A sees 1 red hat and 1 black hat
- All answer “I don’t know” to the first question
- A: If I had red, B would have said “Yes”
- So, after 2<sup>nd</sup> question, A (and B) can say “Yes”

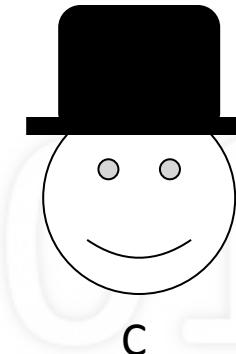
# What If Three Black Hats?



A



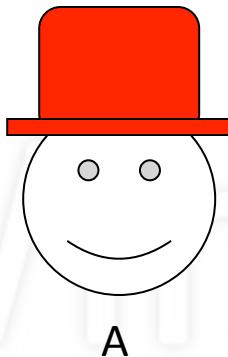
B



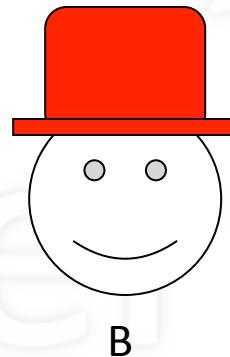
C

- All see 2 black hats
- All answer “I don’t know” to 1<sup>st</sup>, 2<sup>nd</sup> questions
- After 3<sup>rd</sup> question, each can say “Yes”
  - If I had red, others will have said “Yes”
  - Therefore, I must have black

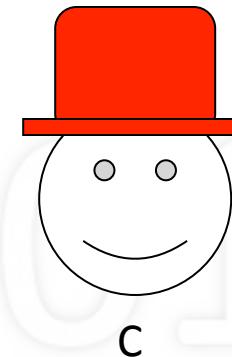
# Why is the Hint Necessary?



A



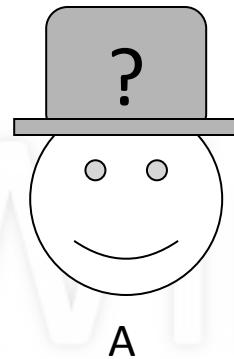
B



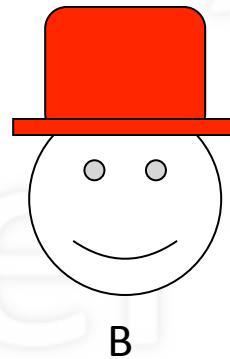
C

- Hint allows each to tell difference between
  - all red hats
  - one black hat
- Provides information about *what others know*

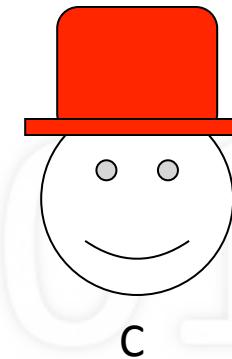
# 1<sup>st</sup> Time: What Does A Know?



A



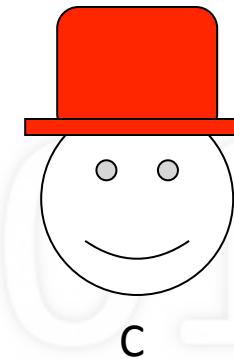
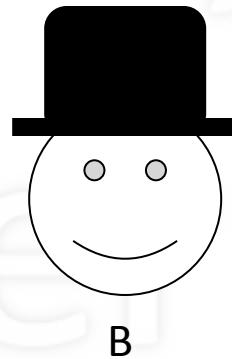
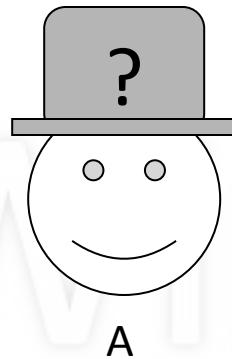
B



C

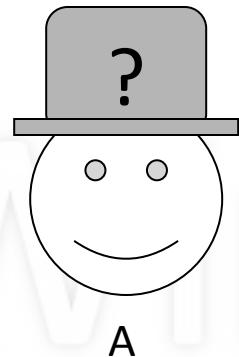
- That B and C have red hats (by observation)
- That there is at least one black hat (the hint)
  - since everyone knows hint
- Deduction: A deduces it must have black hat

# 2<sup>nd</sup> Time: What Does A Know?

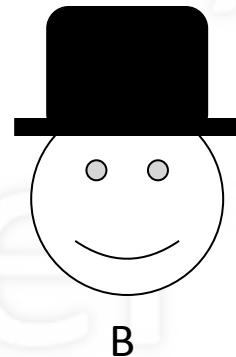


- That B has black, C has red hat (by observation)
- That there is at least one black hat (the hint)
- That B and C know the hint
  - since everyone knows that everyone knows hint
  - Thus, if A had red, B would have said “yes” after 1<sup>st</sup> Q
- Deduction: A deduces it must have black hat

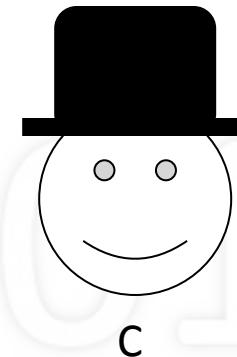
# 3<sup>rd</sup> Time: What Does A Know?



A



B



C

- That B and C have black hats (by observation)
- That there is at least one black hat (the hint)
- That B, C know hint, and that B, C know A knows
  - since everyone knows that everyone knows that everyone knows hint
- Deduction: A deduces it must have black hat!

# Common Knowledge

- Knowledge of x
  - $K(x)$  = Everyone knows x
  - $K^2(x)$ : Everyone knows that everyone knows x
- Common knowledge of x
  - $K(x)$  and  $K^2(x)$  and ...  $K^n(x)$  as  $n \rightarrow \infty$
- Hint is common knowledge
  - Because it is announced publicly
  - Each must assume all heard it and all are smart

# Textbook

- Chapter 19
  - Lecture-related: 19.1, 19.4-19.10
  - Unfortunately, very little material in textbook
  - Try to get a copy of 8<sup>th</sup> edition, or check wikipedia
  - For 8<sup>th</sup> edition (not 9<sup>th</sup>, which also has too little)
    - Lecture-related: 16.1-16.2, 17.1-17.2, Chapter 18 (all)
    - Recommended: 17.3-17.5

# Review & Research

- What is a distributed system?
- What is meant by “degree of integration”?\*
- What are the advantages of a distributed system?
- What is meant by “NSPF” (i.e., no single point of failure)

# R&R

- What are the disadvantages of distributed systems?
- What is meant by “state uncertainty”?
- What is meant by “action uncertainty”?
- Given the advantages and disadvantages, why have distributed systems become popular?\*\*

# R&R

- How do the queueing systems shown in slides 5-6 demonstrate that distribution is better?\*\*
- What is Little's Law (explain the meaning of the parameters)?\*\*\*
- What is the Client/Server Model?\*
- What is the Peer-to-Peer Model?\*
- How do the Client/Server and Peer-to-Peer Models differ?\*\*

# R&R

- What is the purpose of Event Ordering, and how does it work?\*\*
- What is the purpose of Leader Election, and how does the Bully algorithm work?\*\*
- How is mutual exclusion implemented in a distributed system, using the centralized approach vs. the distributed approach?\*\*

# R&R

- What is the Red Hat/Black Hat problem?\*\*
- If there are 3 black hats, why does it take 3 questions before all can say they have a black hat?\*\*
- If there are  $n$  black hats, how many questions are needed before all can say they have a black hat?\*\*\*

# R&R

- Say there are 2 black hats and 1 red hat: how do each figure out they have a black hat?\*\*
- Why is the hint necessary?\*\*\*
- What would happen if the hint were whispered in each person's ear, rather than being told publically?\*\*\*

# R&R

- What is meant by Common Knowledge?\*\*
- What is the relevance of Common Knowledge to distributed systems?\*\*\*